Model Fitting II

Carolina A. de Lima Salge
Assistant Professor
Terry College of Business
University of Georgia

Business Intelligence
Spring 2021

Terry College of Business
UNIVERSITY OF GEORGIA

A tibble: 1,338 x 7

age <dbl>	sex <chr>	bmi <dbl>	children <dbl>	smoker <chr>	region <chr>	charges <dbl>
19	female	27.900	0	yes	southwest	16884.924
18	male	33.770	1	no	southeast	1725.552
28	male	33.000	3	no	southeast	4449.462
33	male	22.705	0	no	northwest	21984.471
32	male	28.880	0	no	northwest	3866.855
31	female	25.740	0	no	southeast	3756.622
46	female	33.440	1	no	southeast	8240.590
37	female	27.740	3	no	northwest	7281.506
37	male	29.830	2	no	northeast	6406.411
60	female	25.840	0	no	northwest	28923.137

1–10 of 1,338 rows   Previous 1 2 3 4 5 6 … 100 Next

```
> postResample(pred = p, obs = charges_test$charges)
        RMSE        Rsquared           MAE
  5808.0045894     0.7989742    4184.9721150
```
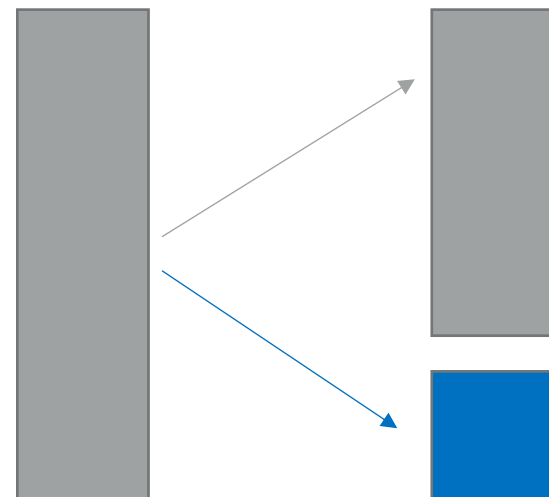
Training Data

Test Data

Data

```
> summary(diff(resamps))

Call:
summary.diff.resamples(object = diff(resamps))

p-value adjustment: bonferroni
Upper diagonal: estimates of the difference
Lower diagonal: p-value for H0: difference = 0

MAE
    LM         RF
LM             1403
RF < 2.2e-16

RMSE
    LM         RF
LM             1156
RF 2.139e-15

Rsquared
    LM         RF
LM            -0.09466
RF 1.17e-14
```
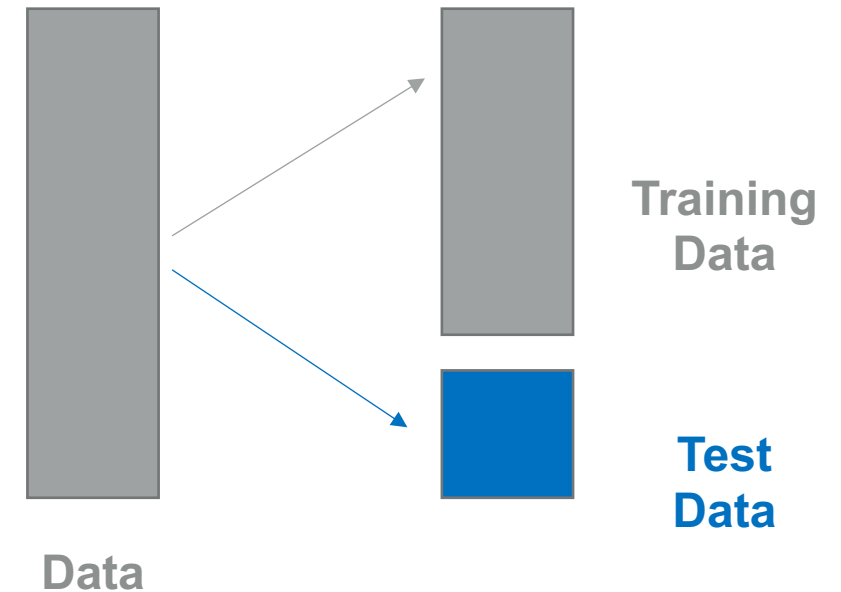
Machine Learning Use

Predictive modeling

- The goal is to predict the target using a new dataset where we have values for predictors but not the target

Machine Learning Use

Evaluate based on prediction error

- # Build model using training data
- # Assess performance on test (hold-out) data

Training
Data

Test
Data

Data

Model Evaluation

How well the model predicts new data (*not* how well it fits the data it was trained with)

- ## Key component of most measures is difference between actual outcome and predicted outcome (i.e., error)

Model Evaluation

Error for data record = predicted (p) minus actual (a)

RMSE: Root Mean Squared Error
MAE: Mean Absolute Error
MAPE: Mean Absolute Percentage Error
Total SSE: Total Sum of Squared Errors

RMSE

Error for data record = predicted (*p*) minus actual (*a*)

RMSE = how much the *p*'s diverge from the *a*'s, on average

Assume the regression equation is y = 1.74x. What is the root mean squared error for the sample dataset?

x	a	p	(p – a)^2
1	2		
2	5		
-1	-2		

RMSE

x	a	p	(p – a)^2
1	2	**1.74**	
2	5	**3.48**	
-1	-2	**-1.74**	

y = 1.74x
        = 1.74 * 1 **= 1.74**
        = 1.74 * 2 **= 3.48**
        = 1.74 * (-1) **= -1.74**

RMSE

x	a	p	(p – a)^2
1	2	1.74	0.0676
2	5	3.48	2.3104
-1	-2	-1.74	0.0676

= (1.74 – 2)^2 = 0.0676
= (3.48 – 5)^2 = 2.3104
= (-1.74 – -2)^2 = 0.0676

RMSE

x	a	p	(p – a)^2
1	2	1.74	0.0676
2	5	3.48	2.3104
-1	-2	-1.74	0.0676

RMSE = (0.0676 + 2.3104 + 0.0676) / 3

$$= \sqrt{0.8152}$$

$$= .903$$

ML Regression in R

Use the the **caret** package

Insurance dataset – recall, *the goal is to predict the target using a new dataset as best as we can*

```
library(tidyverse)
library(caret)

insurance <- read_csv("insurance.csv")
```

Terry College of Business
UNIVERSITY OF GEORGIA

Insurance Data

Medical Cost Personal Datasets
Insurance Forecast by using Linear Regression

Miri Choi • updated 3 years ago (Version 1)

Data   Tasks (2)   Code (478)   Discussion (11)   Activity   Metadata   Download (54 KB)   New Notebook

Usability 8.8   License   Database: Open Database, Contents: Database Contents   Tags education, health, finance, insurance, healthcare

Description

Context

Machine Learning with R by Brett Lantz is a book that provides an introduction to machine learning using R. As far as I can tell, Packt Publishing does not make its datasets available online unless you buy the book and create a user account which can be a problem if you are checking the book out from the library or borrowing the book from a friend. All of these datasets are in the public domain but simply needed some cleaning up and recoding to match the format in the book.

Content

Columns

- age: age of primary beneficiary
- sex: insurance contractor gender, female, male

ML Regression in R

`insurance`

A tibble: 1,338 x 7

age <dbl>	sex <chr>	bmi <dbl>	children <dbl>	smoker <chr>	region <chr>	charges <dbl>
19	female	27.900	0	yes	southwest	16884.924
18	male	33.770	1	no	southeast	1725.552
28	male	33.000	3	no	southeast	4449.462
33	male	22.705	0	no	northwest	21984.471
32	male	28.880	0	no	northwest	3866.855
31	female	25.740	0	no	southeast	3756.622
46	female	33.440	1	no	southeast	8240.590
37	female	27.740	3	no	northwest	7281.506
37	male	29.830	2	no	northeast	6406.411
60	female	25.840	0	no	northwest	28923.137

1–10 of 1,338 rows

Previous 1 2 3 4 5 6 … 100 Next

Selecting Predictors

The goal is to find a parsimonious model – i.e., a simple model that performs well

- Correlation between predictors
- Correlation between predictors and target

Selecting Predictors

To compute the correlation, we need numeric values

```
# transform categories to numbers
insurance <- insurance %>%
  mutate(sexN = case_when(
    sex == "male" ~ 1,
    sex == "female" ~ 0
    )) %>%
  mutate(smokerN = case_when(
    smoker == "yes" ~ 1,
    smoker == "no" ~ 0
    )) %>%
  mutate(regionN = case_when(
    region == "southwest" ~ 1,
    region == "southeast" ~ 2,
    region == "northwest" ~ 3,
    region == "northeast" ~ 4
    ))
```

```
# only select numeric variables
df <- insurance %>%
  dplyr::select(charges, age, sexN, bmi,
children, smokerN, regionN)

# drop missing values NAs
df1 <- drop_na(df)
```

Terry College of Business
UNIVERSITY OF GEORGIA

Model Induction

Set a starting value so that results are reproducible
Split the data into training and testing

```
set.seed(12L) # set a starting seed to be able to get reproducible results


# partition data
trainIndex <- createDataPartition(df1$charges, # target variable
                                  p = 0.8, # percentage that goes to training
                                  list = FALSE, # results will not be in a list
                                  times = 1) # number of partitions to create


charges_train <- df1[trainIndex, ] # data frame for training
charges_test <- df1[-trainIndex, ] # data frame for testing
```

<u>Selecting Predictors</u>

Compute the correlations using the training data

```
# compute correlation between predictors
cor(charges_train[,2:7])
```

```
> cor(charges_train[,2:7])
                    age          sexN          bmi     children      smokerN       regionN
age        1.000000000 -0.008477524  0.1013065882  0.0569025927 -0.033065845  0.002014072
sexN      -0.008477524  1.000000000  0.0504461319  0.0046657032  0.058039957  0.018623882
bmi        0.101306588  0.050446132  1.0000000000 -0.0007692116 -0.014603519 -0.174424109
children   0.056902593  0.004665703 -0.0007692116  1.0000000000  0.007271808 -0.037174622
smokerN   -0.033065845  0.058039957 -0.0146035194  0.0072718079  1.000000000  0.001804169
regionN    0.002014072  0.018623882 -0.1744241089 -0.0371746216  0.001804169  1.000000000
```

Selecting Predictors

Compute the correlations using the training data

```
# compute correlation between predictors and the target
cor(df1[,1:7])
```

```
> cor(charges_train[,1:7])
               charges          age        sexN           bmi      children     smokerN      regionN
charges      1.00000000   0.293517648   0.044956612   0.1924010282   0.0573252106   0.775551461   0.011725009
age          0.29351765   1.000000000  -0.008477524   0.1013065882   0.0569025927  -0.033065845   0.002014072
sexN         0.04495661  -0.008477524   1.000000000   0.0504461319   0.0046657032   0.058039957   0.018623882
bmi          0.19240103   0.101306588   0.050446132   1.0000000000  -0.0007692116  -0.014603519  -0.174424109
children     0.05732521   0.056902593   0.004665703  -0.0007692116   1.0000000000   0.007271808  -0.037174622
smokerN      0.77555146  -0.033065845   0.058039957  -0.0146035194   0.0072718079   1.000000000   0.001804169
regionN      0.01172501   0.002014072   0.018623882  -0.1744241089  -0.0371746216   0.001804169   1.000000000
```

Model Induction and Testing

Use training set to build model, then predict insurance cost using the test set

```
model <- train(charges ~ age + bmi + smokerN,
               data = charges_train, # use training set
               method = "lm") # linear regression


# now predict outcomes in test set
p <- predict(model, charges_test)
```

Model Performance

Use training set to build model, then predict insurance cost using the test set

```
# how did we do? calculate performance across resamples
# RMSE and R-squared
postResample(pred = p, obs = charges_test$charges)
# on average, our prediction is off by $5,808.00
```

```
> postResample(pred = p, obs = charges_test$charges)
        RMSE      Rsquared           MAE
   5808.0045894    0.7989742 4184.9721150
```

Model Performance

How to improve performance? One way is to try and specify a different method

```
model2 <- train(charges ~ age + bmi + smokerN,
                data = charges_train, # use training set
                method = "ranger") # random forest

# now predict outcomes in test set
p1 <- predict(model2, charges_test)

# how did we do? calculate performance across resamples
# RMSE and R-squared
postResample(pred = p1, obs = charges_test$charges)
# on average, our prediction is off by $4,209.91
```

```
> postResample(pred = p1, obs = charges_test$charges)
         RMSE       Rsquared            MAE
  4205.4282707      0.8933816  2436.1432025
```

Which Model?

So many choices!

Linear Regression

```
method = 'lm'
```

Type: Regression

Tuning parameters:

- `intercept` (intercept)

A model-specific variable importance metric is available.

Random Forest

```
method = 'ranger'
```

Type: Classification, Regression

Tuning parameters:

- `mtry` (#Randomly Selected Predictors)
- `splitrule` (Splitting Rule)
- `min.node.size` (Minimal Node Size)

Required packages: `e1071`, `ranger`, `dplyr`

A model-specific variable importance metric is available.

http://topepo.github.io/caret/train-models-by-tag.html#model-tree

7 `train` Models By Tag

The following is a basic list of model types or relevant characteristics. There entires in these lists are arguable. For example: random forests theoretically use feature selection but effectively may not, support vector machines use L2 regularization etc.

Contents

- Accepts Case Weights
- Bagging
- Bayesian Model
- Binary Predictors Only
- Boosting
- Categorical Predictors Only
- Cost Sensitive Learning
- Discriminant Analysis
- Distance Weighted Discrimination
- Ensemble Model
- Feature Extraction
- Feature Selection Wrapper
- Gaussian Process
- Generalized Additive Model
- Generalized Linear Model
- Handle Missing Predictor Data
- Implicit Feature Selection
- Kernel Method
- L1 Regularization
- L2 Regularization
- Linear Classifier
- Linear Regression
- Logic Regression
- Logistic Regression
- Mixture Model
- Model Tree
- Multivariate Adaptive Regression Splines
- Neural Network
- Oblique Tree
- Ordinal Outcomes
- Partial Least Squares
- Patient Rule Induction Method
- Polynomial Model
- Prototype Models
- Quantile Regression
- Radial Basis Function
- Random Forest
- Regularization
- Relevance Vector Machines

Which Model?

6 Available Models

The models below are available in `train`. The code behind these protocols can be obtained using the function `getModelInfo` or by going to the github repository.

Show [238] entries

Search: [   ]

Model	*method* Value	Type	Libraries	Tuning Paramete
Adaptive-Network-Based Fuzzy Inference System	ANFIS	Regression	frbs	num.labels, max.i
Bayesian Regularized Neural Networks	brnn	Regression	brnn	neurons
Bayesian Ridge Regression	bridge	Regression	monomvn	None
Bayesian Ridge Regression (Model Averaged)	blassoAveraged	Regression	monomvn	None
Cubist	cubist	Regression	Cubist	committees, neighbors

http://topepo.github.io/caret/available-models.html

Comparing Between Models

Is one model statistically better than the other?

```
# first collect the resampling results of each model
resamps <- resamples(list(LM = model,
                          RF = model2))
resamps


# then use a simple t-test to evaluate the null
hypothesis that there is no difference
summary(diff(resamps))
```

```
> resamps

Call:
resamples.default(x = list(LM = model, RF = model2))

Models: LM, RF
Number of resamples: 25
Performance metrics: MAE, RMSE, Rsquared
Time estimates for: everything, final model fit
```

```
> summary(diff(resamps))

Call:
summary.diff.resamples(object = diff(resamps))

p-value adjustment: bonferroni
Upper diagonal: estimates of the difference
Lower diagonal: p-value for H0: difference = 0

MAE
       LM        RF
LM               1403
RF < 2.2e-16


RMSE
       LM        RF
LM               1156
RF 2.139e-15


Rsquared
       LM        RF
LM               -0.09466
RF 1.17e-14
```

<u>Summary</u>

- Regression with ML is different than regression with traditional OLS – one is focused on predictions while the other is focused on explanations

- When building a predictive ML model, split data into training and test sets (70-30 or 80-20)

- Always evaluate the performance of a model with the test data, and experiment with different methods to compare the performances of different models

Thank You!